

C++ STL Masterclass (Modern C++17/20)

This is a practical, example-heavy guide to the C++ Standard Template Library (STL): containers, iterators (and a peek at ranges), algorithms, function objects, utilities, and common patterns. Code is modern C++ (C++17+), but you can compile most examples with C++14.

0) What is the STL?

STL is the core of C++'s Standard Library built around **generic programming**: - **Containers** store elements (e.g., `vector`, `map`). - **Iterators** are generalized pointers to traverse containers. - **Algorithms** operate over iterator ranges (e.g., `sort`, `accumulate`). - **Function objects (functors)**, lambdas, and **comparators** customize behavior. - **Allocators** manage memory (rarely customized; know they exist).

Philosophy: Separate data structures (containers) from operations (algorithms) connected by iterators.

1) Containers at a Glance

Sequence Containers

- `array<T, N>` – fixed size, stack allocation semantics, contiguous.
- `vector<T>` – dynamic array, contiguous; **default workhorse**.
- `deque<T>` – double-ended queue, fast push/pop at both ends, not contiguous.
- `list<T>` – doubly linked list, stable iterators, slow random access.
- `forward_list<T>` – singly linked list, minimal overhead, forward iterators only.

Container Adapters

- `stack<T, Container=deque<T>>`
- `queue<T, Container=deque<T>>`
- `priority_queue<T, Container=vector<T>, Compare=less<T>>`

Associative Containers (ordered, tree-based: typically Red-Black Tree)

- `set<T>` / `multiset<T>` – unique / duplicate keys.
- `map<Key, T>` / `multimap<Key, T>` – key→value, unique / duplicate keys.

Unordered Containers (hash table)

- `unordered_set<T>` / `unordered_multiset<T>`
- `unordered_map<Key, T>` / `unordered_multimap<Key, T>`

Utility/String-ish

- `basic_string` (`std::string`, `std::u16string`, ...), `string_view` (non-owning), `bitset`, `span` (C++20, non-owning view of contiguous memory).

2) Big-O Cheat Sheet & When to Choose What

Container	Access	Insert (end)	Insert (middle)	Erase (middle)	Find	Memory	Notes
<code>array</code>	$O(1)$	-	-	-	-	low	fixed size
<code>vector</code>	$O(1)$ amortized	$O(1)$ amortized (end)	$O(n)$	$O(n)$	$O(n)$	low	contiguous; best cache; invalidates on reallocation
<code>deque</code>	$O(1)$	$O(1)$ ends	$O(n)$	$O(n)$	$O(n)$	moderate	fast push_front/back
<code>list</code>	$O(n)$	$O(1)$ with iterator	$O(1)$ with iterator	$O(1)$ with iterator	$O(n)$	high	stable iterators; no random access
<code>forward_list</code>	$O(n)$	$O(1)$ after pos	$O(1)$ after pos	$O(1)$ after pos	$O(n)$	low	singly-linked
<code>set/map</code>	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	moderate	ordered; iterates in order
<code>unordered_set/map</code>	average $O(1)$	avg $O(1)$	avg $O(1)$	avg $O(1)$	avg $O(1)$	can be high	hash quality & load factor matter
<code>priority_queue</code>	top $O(1)$	push $O(\log n)$	-	pop $O(\log n)$	-	low	heap-based

Quick picks - Need random access & speed? → `vector`. - Need order by key? → `map` / `set`. - Need constant-time average lookup by key? → `unordered_map` / `unordered_set`. - Need frequent push/pop at both ends? → `deque`. - Need stable references/iterators across inserts? → `list` (or node-based maps/sets).

3) Iterators (and a peek at Ranges)

Iterator categories (in increasing power): - **Input/Output**: one-pass read/write. - **Forward**: multi-pass single-direction. - **Bidirectional**: `--` allowed (e.g., `list`, `set`). - **Random Access**: `+`, `-`, `[]` (e.g., `vector`, `deque`). - **Contiguous** (C++20): memory contiguous (`vector`, `string`, `array`).

```
std::vector<int> v{1,2,3};
for (auto it = v.begin(); it != v.end(); ++it) {
    *it += 10; // mutate through iterator
}
```

Range-based for (C++11+):

```
for (int &x : v) x *= 2;
```

C++20 ranges (`<ranges>`): algorithm + pipe syntax on ranges

```
#include <ranges>
#include <vector>
#include <algorithm>

auto odds_doubled = std::vector{1,2,3,4,5}
    | std::views::filter([](int x){return x%2;})
    | std::views::transform([](int x){return x*2;});

std::vector<int> out;
std::ranges::copy(odds_doubled, std::back_inserter(out));
```

4) Algorithm Essentials (Header: `<algorithm>`, `<numeric>`, `<ranges>`)

Non-modifying: `all_of`, `any_of`, `none_of`, `for_each`, `count`, `find`, `equal`, `mismatch`.

Modifying: `copy`, `move`, `swap_ranges`, `fill`, `transform`, `generate`, `replace`, `remove`, `unique`.

Partitioning: `partition`, `stable_partition`, `partition_point`.

Sorting: `sort`, `stable_sort`, `partial_sort`, `nth_element`, `is_sorted`, `inplace_merge`.

Binary search family (require sorted range): `lower_bound`, `upper_bound`, `equal_range`, `binary_search`.

Set algorithms (sorted ranges): `set_union`, `set_intersection`, `set_difference`, `set_symmetric_difference`.

Heap: `make_heap`, `push_heap`, `pop_heap`, `sort_heap`.

Permutations: `next_permutation`, `prev_permutation`.

Numeric (`<numeric>`): `accumulate`, `reduce` (C++17), `inner_product`, `partial_sum`, `exclusive_scan` (C++17), `iota`.

Example: Top-K with `nth_element`

```
std::vector<int> v{7,1,4,9,3,8,2,6,5};
size_t k = 3; // smallest 3
std::nth_element(v.begin(), v.begin()+k, v.end());
v.resize(k); // first k elements are the k smallest (unordered within k)
```

Erase-remove idiom

```
std::vector<int> v{1,2,3,2,4};
v.erase(std::remove(v.begin(), v.end(), 2), v.end()); // remove all 2s
```

Transform

```
std::transform(v.begin(), v.end(), v.begin(), [](int x){ return x*x; });
```

5) `vector` Deep Dive

```
std::vector<int> v; // empty
v.reserve(100); // avoid reallocation (performance)
for (int i = 0; i < 10; ++i) v.push_back(i);

v.emplace_back(42); // constructs in place

// iteration
for (const auto &x : v) { /*...*/ }
```

```

// insert/erase mid (O(n))
v.insert(v.begin()+3, 999);
v.erase(v.begin()+5);

// shrink
v.shrink_to_fit();

// data pointer (contiguous)
int *p = v.data(); p[0] = 123;

```

Iterator invalidation (vector) - Inserting may reallocate → **all** iterators/references invalidated. - Erasing a single element invalidates iterators **from erased pos to end**.

Tip: Use indices or `std::size_t` when you need stability across reallocations.

6) deque, list, forward_list

- `deque`: Fast `push_front` and `push_back`; random access allowed, but memory is segmented.
- `list`: Splice without copying elements:

```

std::list<int> a{1,2,3}, b{4,5};
a.splice(a.begin(), b); // moves all from b to front of a, O(1)

```

- `forward_list`: Singly linked; use `insert_after`, `erase_after`.

When to use `list`? Rarely. Prefer `vector` unless you *truly* need stable iterators, frequent middle insert/erase with iterator positions, or splicing between lists.

7) Ordered Maps/Sets (`map`, `set`)

```

std::map<std::string, int> freq;
freq["apple"]++; // inserts default 0, then ++

// Preferred: avoid default insert when not needed
freq.insert({"banana", 3});

// Find
if (auto it = freq.find("apple"); it != freq.end()) {
    it->second += 10;
}

```

```
// Iterate in key order
for (auto &[k,v] : freq) {
    // structured binding, C++17
}

// Bounds
auto it = freq.lower_bound("b"); // first key not less than "b"
```

- `set` / `map` keep keys ordered by comparator (`std::less` by default). Supply custom comparator for custom order or case-insensitive strings.

```
struct CaseLess {
    bool operator()(const std::string& a, const std::string& b) const {
        return std::lexicographical_compare(
            a.begin(), a.end(), b.begin(), b.end(),
            [](char x, char y){ return std::tolower(x) < std::tolower(y); });
    }
};

std::map<std::string, int, CaseLess> m;
```

Duplicates? Use `multimap` / `multiset`. Retrieve all with `equal_range(key)`.

8) Unordered Maps/Sets (Hash Tables)

```
#include <unordered_map>

std::unordered_map<std::string, int> freq;
freq.reserve(1024); // reduce rehashes

freq["cat"]++;
freq.insert({"dog", 2});

if (auto it = freq.find("cat"); it != freq.end()) {
    // avg O(1)
}

// Custom hash for a struct
struct Point {int x,y;};
struct Hash {
    size_t operator()(const Point& p) const noexcept {
        return std::hash<int>()(p.x) * 1315423911u ^ std::hash<int>()(p.y);
    }
};

struct Eq {
```

```

    bool operator==(const Point& a, const Point& b) const noexcept {
        return a.x==b.x && a.y==b.y;
    }
};
std::unordered_set<Point, Hash, Eq> S;

```

Load factor & rehashing: `bucket_count`, `load_factor()`, `rehash(n)`, `reserve(n)` influence performance.

9) Adapters: `stack`, `queue`, `priority_queue`

```

std::stack<int> st; st.push(1); st.top(); st.pop();
std::queue<int> q; q.push(1); q.front(); q.pop();

// Max-heap by default (largest on top)
std::priority_queue<int> pq;
pq.push(5); pq.push(1); pq.push(10);
// Min-heap
std::priority_queue<int, std::vector<int>, std::greater<int>> minpq;

```

10) Strings & `string_view`

```

std::string s = "hello";
s += " world";

// find/replace
auto pos = s.find("lo");
if (pos != std::string::npos) s.replace(pos, 2, "LO");

// string_view: non-owning, cheap slice
std::string text = "abcdef";
std::string_view sv(text);
auto sub = sv.substr(2,3); // "cde"

```

Warning: `string_view` doesn't own data → dangling if source dies.

11) Smart Iteration Utilities

- `std::begin(c)`, `std::end(c)` work for arrays too.

- Insert iterators: `back_inserter(v)`, `front_inserter(dq)`, `inserter(c, it)`.
- Stream iterators for quick IO glue:

```
std::istream_iterator<int> in(std::cin), eof;
std::vector<int> v(in, eof);
std::ostream_iterator<int> out(std::cout, " ");
std::copy(v.begin(), v.end(), out);
```

12) Common Patterns You'll Reuse

12.1 Counting/Frequency Map

```
std::unordered_map<int,int> cnt;
for (int x : v) cnt[x]++;
```

12.2 Sorting with Custom Key (use lambda)

```
struct Item{int id; std::string name; int score;};
std::vector<Item> a;
std::sort(a.begin(), a.end(), [](const Item& A, const Item& B){
    if (A.score != B.score) return A.score > B.score; // desc by score
    return A.name < B.name;                          // tie-break
});
```

12.3 Stable Partition to Move Odds to Front (preserve order)

```
std::stable_partition(v.begin(), v.end(), [](int x){ return x%2; });
```

12.4 Two-sum using `unordered_map`

```
std::pair<int,int> two_sum(const std::vector<int>& a, int target){
    std::unordered_map<int,int> idx; // value -> index
    for (int i = 0; i < (int)a.size(); ++i){
        if (auto it = idx.find(target - a[i]); it != idx.end())
            return {it->second, i};
        idx[a[i]] = i;
    }
    return {-1,-1};
}
```


12.5 Deduplicate & Sort

```
std::sort(v.begin(), v.end());
v.erase(std::unique(v.begin(), v.end()), v.end());
```

13) Iterator Invalidation Rules (must-know)

- `vector`: `push_back` may invalidate *all*; `erase` invalidates from point to end.
- `deque`: inserting/erasing anywhere except ends may invalidate all; `push/pop` at ends may invalidate iterators to ends.
- `list` / `forward_list`: iterators/reference remain valid except for erased elements.
- **Ordered/unordered maps/sets**: `insert` doesn't invalidate iterators; `erase` invalidates only erased elements; `rehash` in unordered containers invalidates all iterators.

14) `emplace` vs `insert` vs `push`

- `push_back(x)` copies/moves `x` into container.
- `emplace_back(args...)` constructs in place → avoids temporary.
- `insert` inserts existing element(s) or range at pos.

Example (avoid constructing `std::pair` twice):

```
std::map<int, std::string> m;
m.emplace(1, "one"); // constructs the pair inside the map
```

15) Comparators & Custom Ordering

- Comparator must define a **strict weak ordering**.
- For `priority_queue` with custom compare (min-heap of pairs by second):

```
using P = std::pair<int, int>;
std::priority_queue<P, std::vector<P>,
    std::function<bool(const P&, const P&>> pq(
        [](const P& a, const P& b){ return a.second > b.second; }));
```

Prefer transparent comparators for heterogeneous lookup (C++14+):

```

struct StrLess {
    using is_transparent = void; // enables lookup by string_view
    bool operator()(std::string_view a, std::string_view b) const { return a <
b; }
};
std::set<std::string, StrLess> S;
S.find("abc"sv); // no temporary std::string

```

16) Hash Customization (Unordered Containers)

- Provide both `Hash` and `KeyEqual` when your key is a struct.
- Combine fields using standard hashes; consider `boost::hash_combine` or a simple mixing constant.

```

struct Key{int a; int b;};
struct KeyHash {
    size_t operator()(const Key& k) const noexcept {
        size_t h1 = std::hash<int>{}(k.a);
        size_t h2 = std::hash<int>{}(k.b);
        return h1 ^ (h2 + 0x9e3779b97f4a7c15ULL + (h1<<6) + (h1>>2));
    }
};
struct KeyEq {
    bool operator()(const Key& x, const Key& y) const noexcept {
        return x.a==y.a && x.b==y.b;
    }
};
std::unordered_map<Key, int, KeyHash, KeyEq> M;

```

17) Numeric & Utility Goodies

```

#include <numeric>
std::vector<int> v{1,2,3,4};
int sum = std::accumulate(v.begin(), v.end(), 0);

#include <bitset>
std::bitset<8> b(0b10110100);
b.flip(0);

#include <tuple>
auto tup = std::make_tuple(1, 2.5, "hi");

```

```

auto [i, d, s] = tup; // structured binding (C++17)

#include <optional>
std::optional<int> maybe;
maybe = 42;
if (maybe) { /*...*/ }

#include <variant>
std::variant<int, std::string> var = 5;
var = std::string("ok");

```

18) Parallel Algorithms (C++17) — `<execution>`

```

#include <execution>
std::sort(std::execution::par, v.begin(), v.end());

```

Policies: `seq`, `par`, `par_unseq`. Use only with safe operations (no data races, iterators valid, etc.).

19) Ranges (C++20) — quick tour

- Headers: `<ranges>`, namespace `std::ranges` / `std::views`.
- Algorithms become range-aware: `std::ranges::sort(vec);`
- Views are lazy, composable. Common views: `filter`, `transform`, `take`, `drop`, `iota`.

```

#include <ranges>
#include <iostream>
for (int x : std::views::iota(1, 10) | std::views::filter([](int x){return
x%3==0;}))
    std::cout << x << ' '; // 3 6 9

```

20) Practical Gotchas & Tips

- Prefer `vector` unless a measurable need says otherwise.
- Pre-reserve for known sizes: `v.reserve(n);`
- Don't keep iterators across operations that may invalidate them.
- Use `auto` + structured bindings for clarity:

```

for (auto &[k, v] : mymap) { /*...*/ }

```

- For `map` / `set`, avoid `operator[]` on `map` when you don't want insertion; use `find`.
- Erase while iterating safely:

```
for (auto it = v.begin(); it != v.end(); ) {
    if (*it % 2 == 0) it = v.erase(it); else ++it;
}
```

- Prefer algorithm + iterator style over manual loops when possible (clearer + fewer bugs).
- Use `span` and `string_view` for non-owning views to avoid copies.

21) Mini-Recipes (copy-paste ready)

Top N largest elements

```
std::nth_element(v.begin(), v.end()-N, v.end());
std::vector<int> topN(v.end()-N, v.end());
std::sort(topN.begin(), topN.end(), std::greater<>());
```

Median of vector

```
auto mid = v.begin() + v.size()/2;
std::nth_element(v.begin(), mid, v.end());
int median = *mid;
```

K-way merge (merge multiple sorted vectors)

```
struct Node{int val, i, j};
auto cmp = [](const Node& a, const Node& b){ return a.val > b.val; };
std::priority_queue<Node, std::vector<Node>, decltype(cmp)> pq(cmp);
```

Group by key (stable)

```
std::stable_sort(a.begin(), a.end(), [](auto &x, auto &y){return x.key <
y.key;});
auto it = a.begin();
while (it != a.end()) {
    auto jt = std::upper_bound(it, a.end(), it->key, [](auto k, auto &obj){return
k < obj.key;});
    // [it, jt) is the group for key it->key
}
```

```
    it = jt;
}
```

22) Complexity Guarantees Snapshot

- `std::sort` average/worst $O(n \log n)$; `stable_sort` worst $O(n \log^2 n)$ (often $O(n \log n)$).
- `unordered_*` average $O(1)$ for find/insert/erase; worst $O(n)$ (bad hash or many collisions).
- `map/set` operations are $O(\log n)$ due to balanced trees.
- `vector::push_back` amortized $O(1)$; reallocation doubles capacity (implementation-dependent growth factor, commonly 1.5–2x).

23) Practice Exercises (with hints)

1. **Frequency of words:** Read N words, print top- k by frequency (desc), ties by lexicographic order. *Hints:* `unordered_map`, then move into `vector<pair<string, int>>`, sort with custom comparator, or use `partial_sort` for top- k .
2. **Interval scheduling:** Given intervals $[l, r)$, select max non-overlapping. *Hints:* sort by end; greedy.
3. **Distinct in sliding window:** Count distinct numbers in each window of size W . *Hints:* `unordered_map` counts, slide with `--` and `++`.
4. **LRU cache:** Implement with `list` + `unordered_map<Key, list::iterator>`.
5. **Autocomplete:** Given dictionary, return words with prefix p . *Hints:* store in `vector`, sort, then use `lower_bound` on prefix ranges.

If you want, we can turn these into template files you can compile and run.

24) Quick Reference (Headers)

- `<vector>`, `<array>`, `<deque>`, `<list>`, `<forward_list>`, `<stack>`, `<queue>`
- `<set>`, `<map>`, `<unordered_set>`, `<unordered_map>`
- `<algorithm>`, `<numeric>`, `<iterator>`, `<functional>`
- `<string>`, `<string_view>`, `` (C++20), `<bitset>`
- `<tuple>`, `<optional>`, `<variant>`
- `<ranges>` (C++20), `<execution>` (C++17)

25) STL Interview-Style Traps You Should Nail

- **Explain erase-remove:** why two calls? (`remove` shifts, returns new logical end; `erase` shrinks container.)
- `map` vs `unordered_map`: order/log- n vs avg $O(1)$, iterates unordered.
- `stable_sort` vs `sort`: stability matters when key ties must preserve input order.

- `lower_bound` / `upper_bound` / `equal_range` : correct usage and invariants.
 - **Why** `vector` **usually beats** `list` **even for many inserts?** Cache locality dominates.
 - **What invalidates what?** Be precise per container.
-

26) Next Steps

- Pick 2–3 exercises above and implement.
 - Convert loops to algorithms (`std::transform`, `std::accumulate`, etc.).
 - Try C++20 ranges on a small project (log filter, CSV processing).
-

Need a printable cheat-sheet PDF or ready-to-run example files? Ask and I'll generate them here.